

```
In [1]: import pandas as pd
```

```
In [2]: data = pd.read_csv("ACME-HappinessSurvey2020.csv")
```

```
In [3]: data.head()
```

```
Out[3]:
```

	Y	X1	X2	X3	X4	X5	X6
0	0	3	3	3	4	2	4
1	0	3	2	3	5	4	3
2	1	5	3	3	3	3	5
3	0	5	4	3	3	3	5
4	0	5	4	3	3	3	5

```
In [4]: data.shape
```

```
Out[4]: (126, 7)
```

```
In [5]: data.Y.value_counts()
```

```
Out[5]: 1    69  
        0    57  
        Name: Y, dtype: int64
```

```
In [6]: data.isna().sum()
```

```
Out[6]: Y      0  
        X1     0  
        X2     0  
        X3     0  
        X4     0  
        X5     0  
        X6     0  
        dtype: int64
```

```
In [7]: data.describe()
```

```
Out[7]:
```

	Y	X1	X2	X3	X4	X5	X6
count	126.000000	126.000000	126.000000	126.000000	126.000000	126.000000	126.000000
mean	0.547619	4.333333	2.531746	3.309524	3.746032	3.650794	4.253968
std	0.499714	0.800000	1.114892	1.023440	0.875776	1.147641	0.809311
min	0.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000
25%	0.000000	4.000000	2.000000	3.000000	3.000000	3.000000	4.000000
50%	1.000000	5.000000	3.000000	3.000000	4.000000	4.000000	4.000000
75%	1.000000	5.000000	3.000000	4.000000	4.000000	4.000000	5.000000
max	1.000000	5.000000	5.000000	5.000000	5.000000	5.000000	5.000000

Goal is to predict if a customer is happy or not based on the answers they give to questions asked

```
In [8]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import StratifiedKFold, cross_validate, tra
in_test_split, cross_val_score, KFold
from sklearn.metrics import roc_curve, auc, classification_report, confu
sion_matrix, precision_score, recall_score, accuracy_score, precision_r
ecall_curve
```

```
In [9]: Y = data['Y']
X = data.drop('Y', axis = 1)
```

```
In [10]: Y.shape, X.shape
```

```
Out[10]: ((126,), (126, 6))
```

```
In [11]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, train_size =
0.7, random_state = 15)
```

```
In [12]: print("Train Data Dimensions : ", X_train.shape)
print("Test Data Dimensions : ", X_test.shape)
```

```
Train Data Dimensions : (88, 6)
Test Data Dimensions : (38, 6)
```

```
In [13]: logreg = LogisticRegression()  
logreg.fit(X_train,Y_train)
```

```
Out[13]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=  
True,  
                             intercept_scaling=1, l1_ratio=None, max_iter=100,  
                             multi_class='auto', n_jobs=None, penalty='l2',  
                             random_state=None, solver='lbfgs', tol=0.0001, verbo  
se=0,  
                             warm_start=False)
```

```
In [14]: Y_pred = logreg.predict(X_test)  
print('Accuracy of logistic regression classifier on test set: {:.2f}'.f  
ormat(logreg.score(X_test, Y_test)))
```

Accuracy of logistic regression classifier on test set: 0.53


```
In [15]: from sklearn.metrics import f1_score  
f1_score(Y_test, Y_pred, average='macro')
```

Out[15]: 0.45192307692307687

```
In [16]: import matplotlib.pyplot as plt
```

```
In [17]: # get importance  
importance = logreg.coef_[0]  
# summarize feature importance  
for i,v in enumerate(importance):  
    print('Feature: %0d, Score: %.5f' % (i,v))  
# plot feature importance  
plt.bar([x for x in range(len(importance))], importance)  
plt.show()
```

Feature: 0, Score: 0.44203
Feature: 1, Score: -0.05557
Feature: 2, Score: 0.07554
Feature: 3, Score: -0.14277
Feature: 4, Score: 0.18195
Feature: 5, Score: 0.16456



Gradient boost

```
In [19]: from sklearn.ensemble import GradientBoostingClassifier
gradient_boost = GradientBoostingClassifier(random_state=1)
gradient_boost.fit(X_train, Y_train)
Y_pred = gradient_boost.predict(X_test)
print('Accuracy of gradient boost classifier on test set: {:.2f}'.format(
    gradient_boost.score(X_test, Y_test)))
```

Accuracy of gradient boost classifier on test set: 0.63

Should the data be divided into 3 sets? training, test and validation set? but dividing into 3 would reduce the data available for training, training set is already small.

```
In [29]: scores = cross_val_score(gradient_boost, X, Y, cv=10)
scores
```

```
Out[29]: array([0.76923077, 0.53846154, 0.46153846, 0.53846154, 0.61538462,
                0.23076923, 0.83333333, 0.75          , 0.33333333, 0.5          ])
```

```
In [30]: print("Accuracy of gradient boosting with 5-fold validation: %0.2f" % (s
cores.mean()))
```

Accuracy of gradient boosting with 5-fold validation: 0.56

cross fold validation is reducing the accuracy, why?

Random Forest

```
In [32]: from sklearn.ensemble import RandomForestClassifier
RF = RandomForestClassifier(random_state=1)
%time RF.fit(X_train, Y_train)
Y_pred = RF.predict(X_test)
print('Accuracy of Random Forest classifier on test set: {:.2f}'.format(
    RF.score(X_test, Y_test)))
```

CPU times: user 143 ms, sys: 2.73 ms, total: 146 ms

Wall time: 149 ms

Accuracy of Random Forest classifier on test set: 0.66

```
In [33]: scores = cross_val_score(RF, X, Y, cv=10)
scores
```

```
Out[33]: array([0.69230769, 0.38461538, 0.46153846, 0.61538462, 0.61538462,
                0.38461538, 0.66666667, 0.75          , 0.33333333, 0.5          ])
```

```
In [43]: print("Accuracy of Randon Forest with 10-fold validation: %0.2f" % (scores.mean()))
```

Accuracy of Randon Forest with 10-fold validation: 0.54

Again, accuracy reduced with cross validation

```
In [36]: RF.feature_importances_
```

```
Out[36]: array([0.1511827 , 0.1772059 , 0.18402833, 0.16528549, 0.18015894,
                0.14213864])
```

```
In [42]: print("Features sorted by their score:")
print(sorted(zip(map(lambda x: round(x, 4), RF.feature_importances_), X.columns),
              reverse=True))
```

Features sorted by their score:
[(0.184, 'X3'), (0.1802, 'X5'), (0.1772, 'X2'), (0.1653, 'X4'), (0.1512, 'X1'), (0.1421, 'X6')]